

# recruitML

June 7, 2024

## 0.1 : ML

```
[1]: !pip install diffprivlib
```

```
Requirement already satisfied: diffprivlib in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (0.6.4)
Requirement already satisfied: numpy>=1.21.6 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from diffprivlib)
(1.26.4)
Requirement already satisfied: scikit-learn>=0.24.2 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from diffprivlib)
(1.4.2)
Requirement already satisfied: scipy>=1.7.3 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from diffprivlib)
(1.13.0)
Requirement already satisfied: joblib>=0.16.0 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from diffprivlib)
(1.4.2)
Requirement already satisfied: setuptools>=49.0.0 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from diffprivlib)
(69.5.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages (from scikit-
learn>=0.24.2->diffprivlib) (3.5.0)
```

```
[2]: import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB as SklearnGaussianNB
import numpy as np
import spacy
from typing import List, Tuple
import itertools
import matplotlib.pyplot as plt
from diffprivlib.models import GaussianNB as DPGaussianNB
import matplotlib.font_manager as fm

#
```

```

data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)
font_path = '/usr/share/fonts/truetype/fonts-japanese-gothic.ttf' #
font_prop = fm.FontProperties(fname=font_path)
# spaCy
nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

# BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
    ↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
    ↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
    ↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,

```

```

        "negative": 0,
    }
    df["sentiment"] = df["sentiment"].map(m)
    df["bow"] = bow.tolist()

    X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
        ↪test_size=0.2)
    X_train = [list(x) for x in X_train]
    X_test = [list(x) for x in X_test]

    #
    clf = SklearnGaussianNB()
    clf.fit(X_train, y_train)
    print("Non-DP accuracy: ", clf.score(X_test, y_test))

    #
    epsilons = np.logspace(-2, 2, 50)
    dim = np.array(X_train).shape[1]
    lowers = np.zeros(dim)
    uppers = np.ones(dim)
    accuracies = {}

    for epsilon in epsilons:
        accuracy = []
        for _ in range(20):
            dp_clf = DPGaussianNB(bounds=(lowers, uppers), epsilon=epsilon)
            dp_clf.fit(X_train, y_train)
            accuracy.append(dp_clf.score(X_test, y_test))
        accuracies[epsilon] = accuracy

    #
    x = epsilons
    y = [np.mean(accuracies[eps]) for eps in epsilons]
    e = [np.std(accuracies[eps]) for eps in epsilons]

    plt.figure(figsize=(10, 6))
    plt.semilogx(x, y)
    plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
    plt.xlabel(' ', fontproperties=font_prop)
    plt.ylabel('accuracy')
    plt.ylim(0, 1)
    plt.title('          accuracy ', fontproperties=font_prop)
    plt.grid(True)
    plt.show()


```

/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-packages/torch/cuda/\_\_init\_\_.py:118: UserWarning: CUDA initialization: CUDA

unknown error - this may be due to an incorrectly set up environment, e.g. changing env variable CUDA\_VISIBLE\_DEVICES after program start. Setting the available devices to be zero. (Triggered internally at `../c10/cuda/CUDAFunctions.cpp:108.`)

```
return torch._C._cuda_getDeviceCount() > 0
/tmp/ipykernel_948393/3754234034.py:46: UserWarning: [W008] Evaluating
Token.similarity based on empty vectors.
similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
parsed_doc if token.pos_ in POS]
```

Non-DP accuracy: 0.6444644464446445



## 0.2

```
[3]: from sklearn.linear_model import LogisticRegression
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB as SklearnGaussianNB
import numpy as np
import spacy
from typing import List, Tuple
import itertools
import matplotlib.pyplot as plt
```

```

from difflib import GaussianNB as DPGaussianNB
import matplotlib.font_manager as fm

#
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)
font_path = '/usr/share/fonts/truetype/fonts-japanese-gothic.ttf' #
font_prop = fm.FontProperties(fname=font_path)
# spaCy
nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

# BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
    ↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
    ↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
    ↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

```

```

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = bow.tolist()

X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪test_size=0.2)
X_train = [list(x) for x in X_train]
X_test = [list(x) for x in X_test]
#
clf = LogisticRegression(random_state=0).fit(X_train, y_train.to_numpy())
print("Non-DP accuracy: ", clf.score(X_test, y_test.to_numpy()))

```

/tmp/ipykernel\_948393/1424478305.py:47: UserWarning: [W008] Evaluating  
Token.similarity based on empty vectors.  
similarities = [(token.similarity(parsed\_doc), token.lemma\_) for token in  
parsed\_doc if token.pos\_ in POS]

Non-DP accuracy: 0.7155715571557155

```

[4]: # from sklearn.linear_model import LogisticRegression
# clf = LogisticRegression(random_state=0).fit(X_train, y_train.to_numpy())
# print("accuracy: ", clf.score(X_test, y_test.to_numpy()))
import math
import numpy as np
import matplotlib.pyplot as plt
from difflib.models import LogisticRegression as DPLR

epsilons = np.logspace(-2, 2, 50)
dim = np.array(X_train).shape[1]
data_norm = math.sqrt(dim)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=data_norm, epsilon=epsilon).fit(X_train, y_train.
    ↪to_numpy())
        accuracy.append(clf.score(X_test, y_test.to_numpy()))
    accuracies[epsilon] = accuracy

#
x = epsilons

```


```

y = [np.mean(accuracies[eps]) for eps in epsilons]
e = [np.std(accuracies[eps]) for eps in epsilons]

plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsizes=10, lw=1)
plt.xlabel(' ')
plt.ylabel('accuracy')
plt.ylim(0, 1)
plt.title('          accuracy ', fontproperties=font_prop)
plt.grid(True)
plt.show()

# import math
# import numpy as np
# import matplotlib.pyplot as plt
# from diffprivlib.models import LogisticRegression
# epsilons = np.logspace(-2, 2, 50)
# dim = np.array(X_train).shape[1]
# data_norm = math.sqrt(dim)
# accuracies = {}
# for epsilon in epsilons:
#     accuracy = []
#     for i in range(20):
#         clf = LogisticRegression(data_norm=data_norm, epsilon=epsilon).
#         ↪fit(X_train, y_train.to_numpy())
#         accuracy.append(clf.score(X_test, y_test.to_numpy()))
#     accuracies[epsilon] = accuracy

```



### 0.3

```
[3]: from sklearn.linear_model import LogisticRegression
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB as SklearnGaussianNB
import numpy as np
import spacy
from typing import List, Tuple
import itertools
import matplotlib.pyplot as plt
from diffprivlib.models import LogisticRegression as DPLR
import matplotlib.font_manager as fm

#
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)
font_path = '/usr/share/fonts/truetype/fonts-japanese-gothic.ttf' #
font_prop = fm.FontProperties(fname=font_path)

# spaCy
nlp = spacy.load('ja_ginza')
```



```

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

# BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
        ↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
        ↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
    ↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = bow.tolist()

```

```

X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪test_size=0.2)
X_train = np.array(X_train.tolist())
X_test = np.array(X_test.tolist())

#
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

#
clf = LogisticRegression(random_state=0).fit(X_train_scaled, y_train.to_numpy())
print("Non-DP accuracy: ", clf.score(X_test_scaled, y_test.to_numpy()))

#
import math

epsilons = np.logspace(-2, 2, 50)
dim = np.array(X_train_scaled).shape[1]
data_norm = math.sqrt(dim)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=data_norm, epsilon=epsilon).fit(X_train_scaled,
    ↪y_train.to_numpy())
        accuracy.append(clf.score(X_test_scaled, y_test.to_numpy()))
    accuracies[epsilon] = accuracy

#
x = epsilons
y = [np.mean(accuracies[eps]) for eps in epsilons]
e = [np.std(accuracies[eps]) for eps in epsilons]


plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
plt.xlabel(' ', fontproperties=font_prop)
plt.ylabel('accuracy', fontproperties=font_prop)
plt.ylim(0, 1)
plt.title('          accuracy ', fontproperties=font_prop)
plt.grid(True)
plt.show()


```

/tmp/ipykernel\_984731/2882912114.py:49: UserWarning: [W008] Evaluating Token.similarity based on empty vectors.

```
similarities = [(token.similarity(parsed_doc), token.lemma_) for token in  
parsed_doc if token.pos_ in POS]
```

Non-DP accuracy: 0.7020702070207021





0.4

```
[4]: from sklearn.linear_model import LogisticRegression
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np
import spacy
from typing import List, Tuple
import itertools
import matplotlib.pyplot as plt
from diffprivlib.models import LogisticRegression as DPLR
import matplotlib.font_manager as fm

#
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)
font_path = '/usr/share/fonts/truetype/fonts-japanese-gothic.ttf' #
font_prop = fm.FontProperties(fname=font_path)

# spacy
```

```

nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

#    BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = bow.tolist()

```

```

X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪test_size=0.2)
X_train = np.array(X_train.tolist())
X_test = np.array(X_test.tolist())

#
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

#
clf = LogisticRegression(random_state=0).fit(X_train_scaled, y_train.to_numpy())
print("Non-DP accuracy: ", clf.score(X_test_scaled, y_test.to_numpy()))

#
import math

#
epsilons = np.logspace(-2, 3, 100)
dim = np.array(X_train_scaled).shape[1]
data_norm = math.sqrt(dim)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=data_norm, epsilon=epsilon).fit(X_train_scaled,
    ↪y_train.to_numpy())
        accuracy.append(clf.score(X_test_scaled, y_test.to_numpy()))
    accuracies[epsilon] = accuracy

#
x = epsilons
y = [np.mean(accuracies[eps]) for eps in epsilons]
e = [np.std(accuracies[eps]) for eps in epsilons]

plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
plt.xlabel(' ', fontproperties=font_prop)
plt.ylabel('accuracy', fontproperties=font_prop)
plt.ylim(0, 1)
plt.title('          accuracy ', fontproperties=font_prop)
plt.grid(True)
plt.show()

```

/tmp/ipykernel\_984731/728075650.py:48: UserWarning: [W008] Evaluating

Token.similarity based on empty vectors.

```
similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
parsed_doc if token.pos_ in POS]
```

Non-DP accuracy: 0.7083708370837084

```
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
  warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
```

```

warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-

```

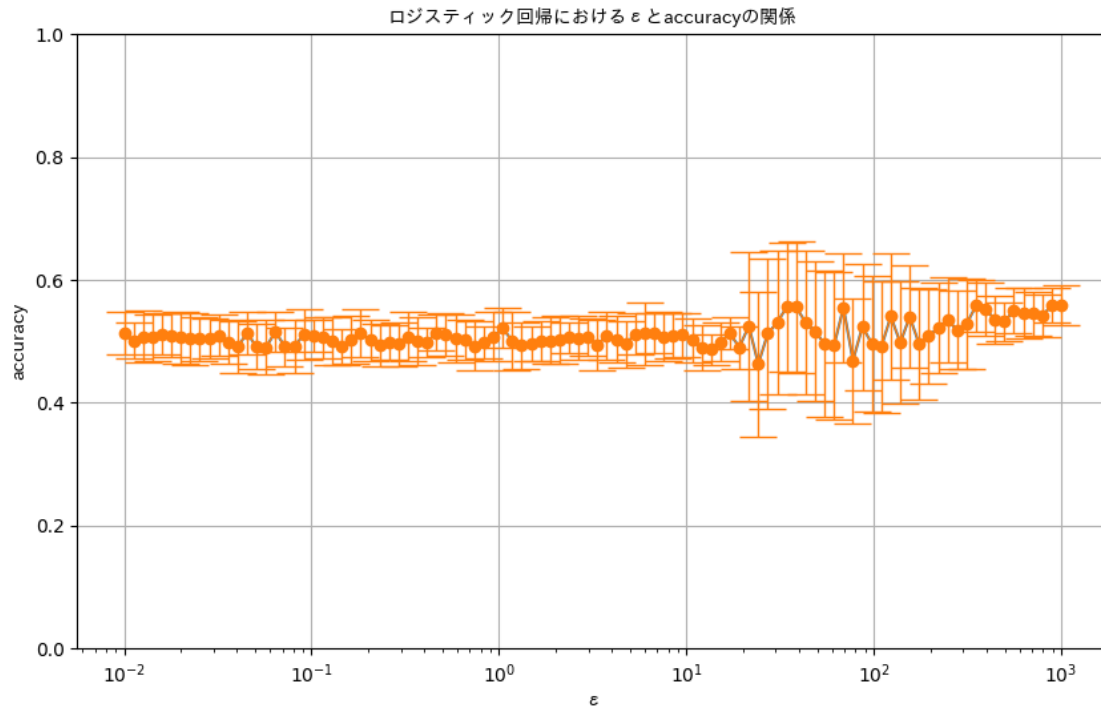




```

ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/diffprivlib/models/logistic_regression.py:423: ConvergenceWarning:
lbfgs failed to converge. Increase the number of iterations.
    warnings.warn("lbfgs failed to converge. Increase the number of iterations.",
ConvergenceWarning)

```



## 0.5 TF-IDF

```
[5]: #
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)

#
print(df.isnull().sum())

#
df = df.dropna()

# spaCy
nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
```

```

res = []
for l in list(itertools.chain.from_iterable(lists)):
    for e in l:
        res.append(e)
return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

# BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
    ↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
    ↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
    ↪ max_features=NUM_VOCAB)
tfidf_bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = tfidf_bow.tolist()

#
print(df.isnull().sum())

#
df = df.dropna()

X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪ test_size=0.2)
X_train = [list(x) for x in X_train]
X_test = [list(x) for x in X_test]

```

```
#
clf = LogisticRegression(random_state=0).fit(X_train, y_train.to_numpy())
print("Non-DP accuracy: ", clf.score(X_test, y_test.to_numpy()))
```

```
review      0
sentiment   0
dtype: int64
```

```
/tmp/ipykernel_984731/346474164.py:39: UserWarning: [W008] Evaluating
Token.similarity based on empty vectors.
```

```
similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
parsed_doc if token.pos_ in POS]
```

```
review      0
sentiment   0
bow         0
dtype: int64
```

```
Non-DP accuracy:  0.7155715571557155
```

```
[6]: import math
import numpy as np
import matplotlib.pyplot as plt
from diffprivlib.models import LogisticRegression as DPLR

epsilons = np.logspace(-2, 2, 50)
dim = np.array(X_train).shape[1]
data_norm = math.sqrt(dim)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=data_norm, epsilon=epsilon).fit(X_train, y_train.
→to_numpy())
        accuracy.append(clf.score(X_test, y_test.to_numpy()))
    accuracies[epsilon] = accuracy

#
x = epsilons
y = [np.mean(accuracies[eps]) for eps in epsilons]
e = [np.std(accuracies[eps]) for eps in epsilons]

plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
plt.xlabel(' ')
plt.ylabel('accuracy')
plt.ylim(0, 1)
```

```
plt.title('          accuracy  (TF-IDF )')
plt.grid(True)
plt.show()
```

```
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12525 (\N{KATAKANA
LETTER RO}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12472 (\N{KATAKANA
LETTER ZI}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12473 (\N{KATAKANA
LETTER SU}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12486 (\N{KATAKANA
LETTER TE}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12451 (\N{KATAKANA
LETTER SMALL I}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12483 (\N{KATAKANA
LETTER SMALL TU}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12463 (\N{KATAKANA
LETTER KU}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 22238 (\N{CJK
UNIFIED IDEOGRAPH-56DE}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 24112 (\N{CJK
UNIFIED IDEOGRAPH-5E30}) missing from current font.
```

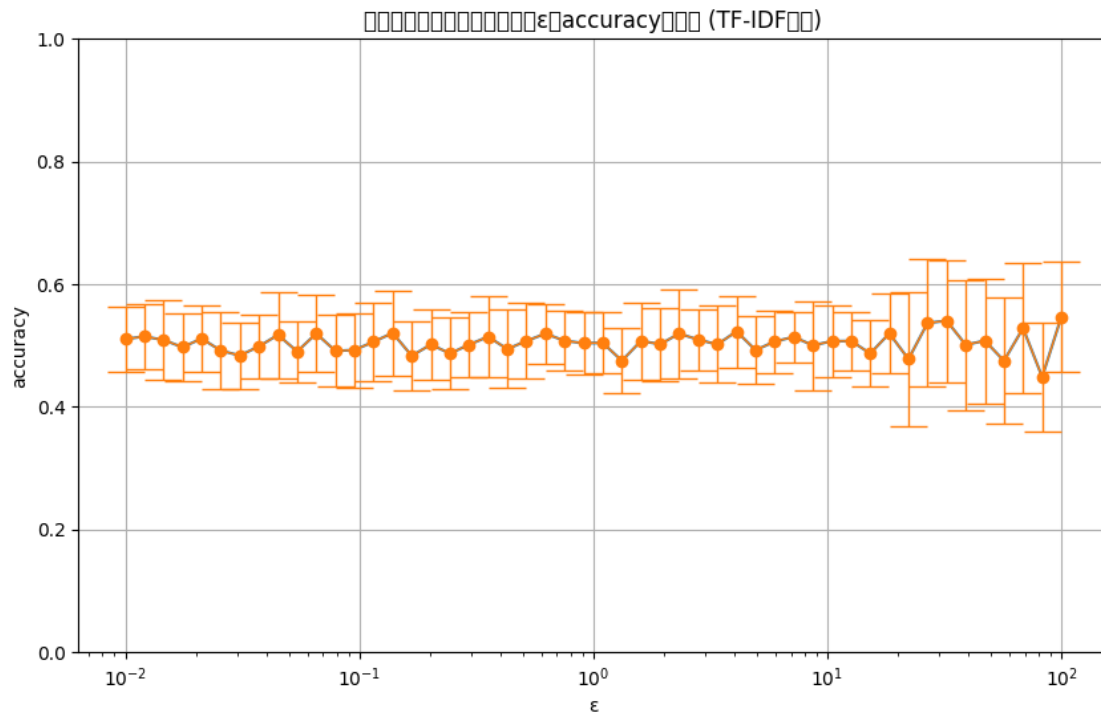
```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12395 (\N{HIRAGANA
LETTER NI}) missing from current font.
```

```
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12362 (\N{HIRAGANA
LETTER O}) missing from current font.
```

```

fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12369 (\N{HIRAGANA
LETTER KE}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12427 (\N{HIRAGANA
LETTER RU}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12392 (\N{HIRAGANA
LETTER TO}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 12398 (\N{HIRAGANA
LETTER NO}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 38306 (\N{CJK
UNIFIED IDEOGRAPH-95A2}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 20418 (\N{CJK
UNIFIED IDEOGRAPH-4FC2}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 20351 (\N{CJK
UNIFIED IDEOGRAPH-4F7F}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)
/home/jun/.pyenv/versions/3.11.8/lib/python3.11/site-
packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 29992 (\N{CJK
UNIFIED IDEOGRAPH-7528}) missing from current font.
fig.canvas.print_figure(bytes_io, **kw)

```



0.6       $\Lambda$       +

```
[7]: import math
import numpy as np
import pandas as pd
import spacy
from typing import List, Tuple
from sklearn.linear_model import LogisticRegression as SklearnLogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import GaussianNB as SklearnGaussianNB
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from diffprivlib.models import LogisticRegression as DPLR
import matplotlib.font_manager as fm

#
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)
font_path = '/usr/share/fonts/truetype/fonts-japanese-gothic.ttf' #
font_prop = fm.FontProperties(fname=font_path)

# spaCy
```



```

nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

#    BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = bow.tolist()

```

```

X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪test_size=0.2)
X_train = np.array(X_train.tolist())
X_test = np.array(X_test.tolist())

#
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

#
def clamp(data, lower, upper):
    ↪return np.clip(data, lower, upper)

#
lower_bound = -100
upper_bound = 100

#
plt.figure(figsize=(10, 6))
plt.hist(X_train_scaled.flatten(), bins=50, alpha=0.5, label='Before Clamping')
plt.xlabel('Value', fontproperties=font_prop)
plt.ylabel('Frequency', fontproperties=font_prop)
plt.title('      (      )', fontproperties=font_prop)
plt.legend()
plt.grid(True)
plt.show()

X_train_clamped = clamp(X_train_scaled, lower_bound, upper_bound)
X_test_clamped = clamp(X_test_scaled, lower_bound, upper_bound)

#
plt.figure(figsize=(10, 6))
plt.hist(X_train_clamped.flatten(), bins=50, alpha=0.5, label='After Clamping',
    ↪color='orange')
plt.xlabel('Value', fontproperties=font_prop)
plt.ylabel('Frequency', fontproperties=font_prop)
plt.title('      (      )', fontproperties=font_prop)
plt.legend()
plt.grid(True)
plt.show()

#
clf = SklearnLogisticRegression(random_state=0).fit(X_train_clamped, y_train.
    ↪to_numpy())
print("Non-DP accuracy: ", clf.score(X_test_clamped, y_test.to_numpy()))

```

```

#
epsilons = np.logspace(-2, 2, 50)
dim = np.array(X_train_clamped).shape[1]
data_norm = math.sqrt(dim)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=data_norm, epsilon=epsilon).fit(X_train_clamped,
↪y_train.to_numpy())
        accuracy.append(clf.score(X_test_clamped, y_test.to_numpy()))
    accuracies[epsilon] = accuracy

#
def bias_correction(accuracies, epsilons):
    corrected_accuracies = {}
    for epsilon in epsilons:
        corrected_accuracies[epsilon] = []
        for acc in accuracies[epsilon]:
            corrected_accuracies[epsilon].append(acc * (1 + (1 / epsilon)))
    return corrected_accuracies

corrected_accuracies = bias_correction(accuracies, epsilons)

#
x = epsilons
# y = [np.mean(corrected_accuracies[eps]) for eps in epsilons]
# e = [np.std(corrected_accuracies[eps]) for eps in epsilons]
y = [np.mean(accuracies[eps]) for eps in epsilons]
e = [np.std(accuracies[eps]) for eps in epsilons]

plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
plt.xlabel(' ', fontproperties=font_prop)
plt.ylabel('accuracy', fontproperties=font_prop)
plt.ylim(0, 1)
plt.title('          accuracy ', fontproperties=font_prop)
plt.grid(True)
plt.show()



```

/tmp/ipykernel\_984731/200654788.py:49: UserWarning: [W008] Evaluating Token.similarity based on empty vectors.


```

similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
parsed_doc if token.pos_ in POS]

```



Non-DP accuracy: 0.7065706570657065



## 0.7

```
[24]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression as SklearnLR
from diffprivlib.models import LogisticRegression as DPLR
import spacy
from typing import List, Tuple
import itertools
import math

#
data_path = './data/reviews_with_sentiment.csv'
df = pd.read_csv(data_path)

# spaCy
nlp = spacy.load('ja_ginza')

#
POS = ['ADJ', 'ADV', 'INTJ', 'PROPN', 'NOUN', 'VERB']
MAX_TERMS_IN_DOC = 5
```

```

NGRAM = 1
MAX_DF = 1.0
MIN_DF = 0.01
NUM_VOCAB = 10000

def flatten(*lists) -> list:
    res = []
    for l in list(itertools.chain.from_iterable(lists)):
        for e in l:
            res.append(e)
    return res

def remove_duplicates(l: List[Tuple[str, float]]) -> List[Tuple[str, float]]:
    d = {}
    for e in l:
        d[e[0]] = e[1]
    return list(d.items())

# BoW
tokens = []
for doc in df["review"]:
    parsed_doc = nlp(doc)
    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
    ↪ parsed_doc if token.pos_ in POS]
    similarities = remove_duplicates(similarities)
    similarities = sorted(similarities, key=lambda sim: sim[1], reverse=True)[:
    ↪ MAX_TERMS_IN_DOC]
    tokens.append([similarity[1] for similarity in similarities])

cv = CountVectorizer(ngram_range=(1, NGRAM), max_df=MAX_DF, min_df=MIN_DF,
    ↪ max_features=NUM_VOCAB)
bow = cv.fit_transform([" ".join(ts) for ts in tokens]).toarray()

#
m = {
    "positive": 1,
    "neutral": 0,
    "negative": 0,
}
df["sentiment"] = df["sentiment"].map(m)
df["bow"] = bow.tolist()
df

```

/tmp/ipykernel\_984731/1391980650.py:45: UserWarning: [W008] Evaluating Token.similarity based on empty vectors.

```

    similarities = [(token.similarity(parsed_doc), token.lemma_) for token in
    parsed_doc if token.pos_ in POS]

```

```
[24]:
```

	review	sentiment \
0		0
1		0
2		1
3		1
4		1
...	...	...
5548	0	
5549		0
5550	0	
5551	0	
5552	...	0

	bow
0	[0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, ...
1	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
2	[0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, ...
3	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, ...
4	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, ...
...	...
5548	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
5549	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
5550	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
5551	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
5552	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...

[5553 rows x 3 columns]

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(df["bow"], df["sentiment"],
    ↪test_size=0.2)
X_train = np.array(X_train.tolist())
X_test = np.array(X_test.tolist())

#
lower_bound = 0
upper_bound = 1.0
clamped_values = np.clip(X_train, lower_bound, upper_bound)

#
plt.figure(figsize=(10, 6))
plt.hist(X_train.flatten(), bins=50, alpha=0.5, label='Before Clamping')
plt.hist(clamped_values.flatten(), bins=50, alpha=0.5, label='After Clamping')
plt.legend()
plt.title('')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.grid(True)
```

```

plt.show()

#
clf = SklearnLR(random_state=0).fit(X_train, y_train)
print("Non-DP accuracy: ", clf.score(X_test, y_test))

#
epsilons = np.logspace(-2, 2, 50)
accuracies = {}

for epsilon in epsilons:
    accuracy = []
    for i in range(20):
        clf = DPLR(data_norm=upper_bound, epsilon=epsilon).fit(clamped_values,
↪y_train)
        accuracy.append(clf.score(X_test, y_test))
    accuracies[epsilon] = accuracy

#
def bias_correction(estimated_values, clamped_values, lower, upper):
    alpha1 = np.mean(estimated_values < lower)
    alpha2 = np.mean(estimated_values > upper)
    theta_T = np.mean(clamped_values[(estimated_values >= lower) &
↪(estimated_values <= upper)])
    corrected_estimates = -alpha1 * lower + (1 - alpha2 - alpha1) * theta_T +
↪alpha2 * upper
    return corrected_estimates

corrected_accuracies = {}
for eps in epsilons:
    estimated_values_for_eps = np.array(accuracies[eps])
    clamped_estimated_values = np.clip(estimated_values_for_eps, lower_bound,
↪upper_bound)
    corrected_accuracy = bias_correction(clamped_estimated_values,
↪clamped_values.flatten(), lower_bound, upper_bound)
    corrected_accuracies[eps] = corrected_accuracy

#
x = epsilons
y = [np.mean(corrected_accuracies[eps]) for eps in epsilons]
e = [np.std(corrected_accuracies[eps]) for eps in epsilons]

plt.figure(figsize=(10, 6))
plt.semilogx(x, y)
plt.errorbar(x, y, yerr=e, marker='o', capthick=1, capsize=10, lw=1)
plt.xlabel(' ')
plt.ylabel('accuracy')

```



```
plt.ylim(0, 1)
plt.title('          accuracy ', fontproperties=font_prop)
plt.grid(True)
plt.show()
```